


Intro to Teradata

Data Warehousing


Amount of Detail Data



Number of Concurrent Users



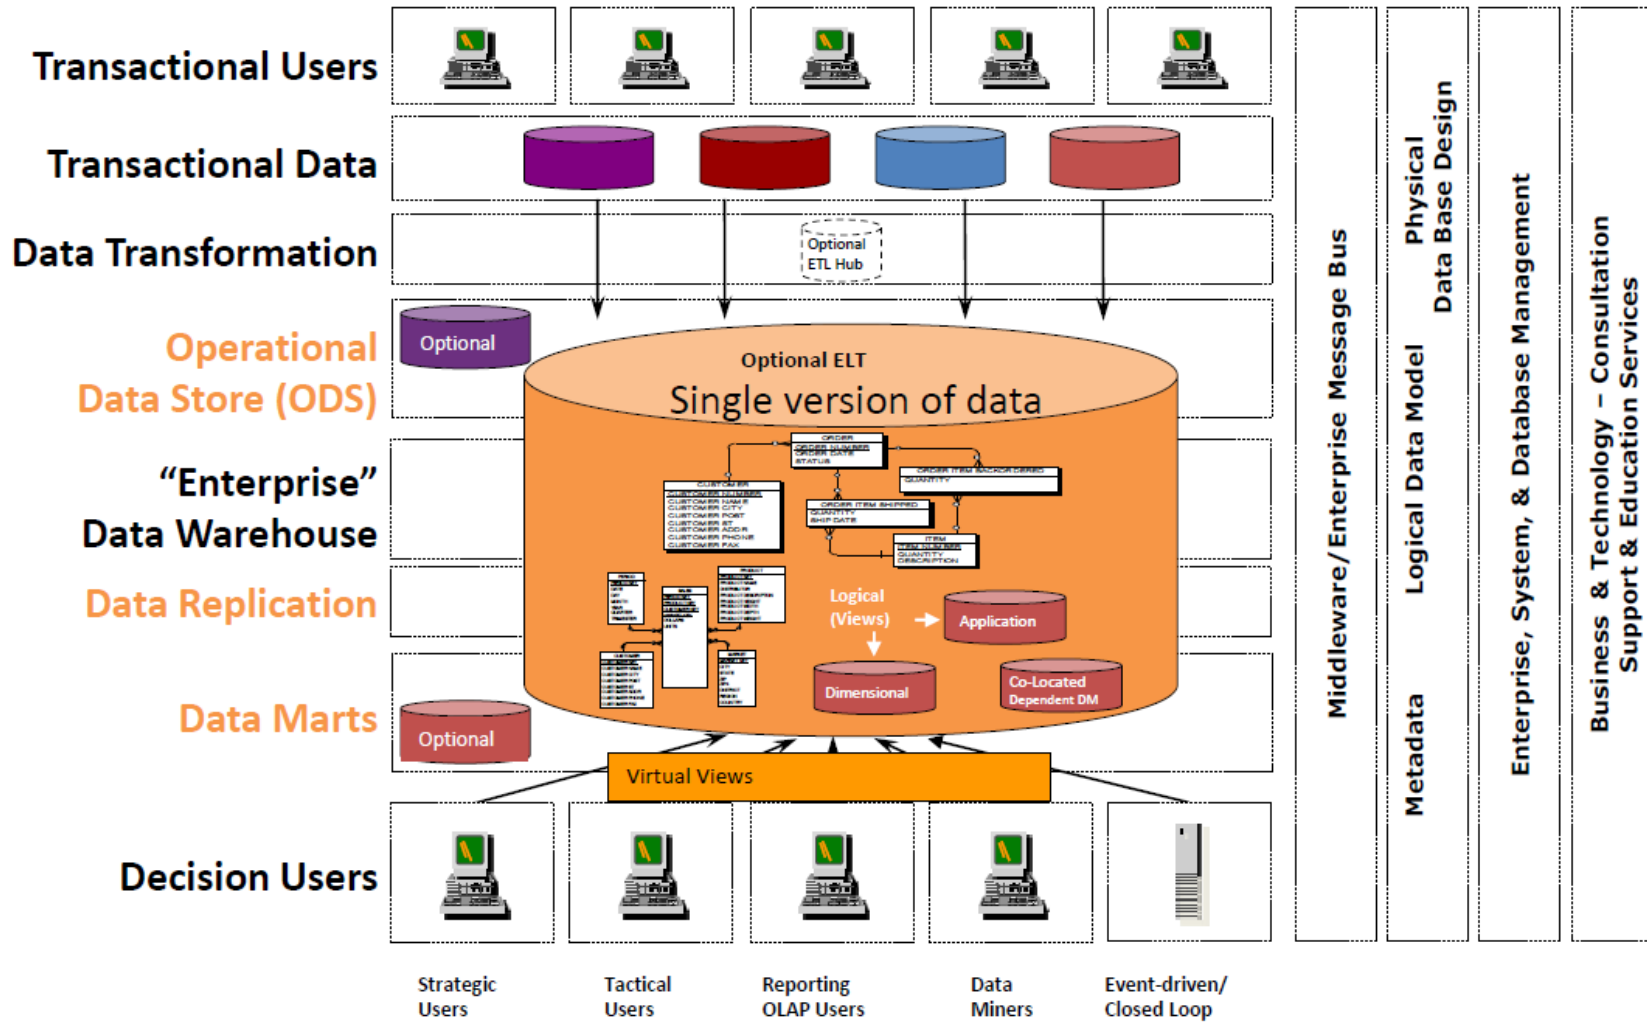
Complexity of the Data Model



Query Complexity

- Simple, direct index
- Moderate, multitable join
- Complex, 10-way table join; includes regression analysis

RDBMS



What is Teradata?

- A global leader in enterprise data warehousing and analytic technologies.
- In data warehousing arena since 1979 in over 60 countries delivering data warehouse appliance solutions.
- Has about 1,000 customers, with annual revenue of about \$1.6 billion from data warehousing solutions.
- Provides functionality for large, complex, mixed workload environments.
- Newer products (Teradata Vantage) that support relational and non-relational data.

What is Teradata (cont.)?

- It has specific strengths (e.g., strong penetration, data models and professional services) in the vertical markets such as retail, financial and banking, telecom and manufacturing.
- Its largest and most prominent customer is Wal-Mart. Other customers include companies like O2, FedEx and eBay.
- Main competitors are other mature DBMS solutions such as IBM's Netezza, Microsoft SQL Server and Oracle.
- Latest data warehouse release is Teradata 16.

What are Teradata features?

- An RDBMS designed for enterprise data warehousing.
- A massively parallel processing system to ensure that tasks are processed quickly.
- A **shared nothing** architecture.
- **Linear scalability** in all dimensions of a database system workload (i.e., data volume, breadth, number of users, complexity of queries). Can scale up to 2048 nodes.


What are Teradata features (cont.)?

- **Connectivity:** to channel-attached systems such as Mainframe or Network-attached systems.
- **Mature Optimizer:** Quite good from the start, it has been refined for each release.
- **SQL:** Industry standard SQL. In addition to this, it provides its own extension.
- **Robust Utilities:** for import/export data from/to such as `FastLoad`, `MultiLoad`, `FastExport` and `TPT`.
- **Automatic Distribution:** No manual intervention for data redistribution.

Parallelism, the secret sauce

Performance & Capacity & Scalability

Break Apart the Query...




Parallel Architectures

Shared Everything

well known RDBMS

- A single database buffer used by all UoP's
- A single logical data store accessed by all UoP's
- Scalability limited due to control bottlenecks and scalability of single SMP platform




◆ - Unit of Parallelism

Shared Nothing

Teradata

- Each UoP is assigned a data portion
- Query Controller ships functions to UoP's that own the data
- Locks, buffers, etc. not shared
- Highly scalable data volumes



Teradata
a division of NCR

Teradata Architecture

- Based on Massively Parallel Processing (MPP) architecture.
- **Massively parallel processing (MPP)**: Multiple SMP nodes working together comprise a larger, MPP implementation. The nodes are connected using a message passing layer called BYNET, which allows multiple virtual processors on multiple nodes to communicate with each other.
- The major components of Teradata are **Parsing Engine, BYNET** and Access Module Processors (**AMPs**).

Components of Teradata

The key components of Teradata are:

- **Node:** It is the basic unit in Teradata System. Each individual server in a Teradata system is referred as a Node.
- A node consists of its own operating system, CPU, memory, own copy of Teradata RDBMS software and disk space.
- A **cabinet** consists of one or more Nodes.

Components of Teradata (cont.)

- **Parsing Engine:** Parsing Engine is responsible for receiving queries from the client and preparing an efficient execution plan.
- The responsibilities of parsing engine are:
 - Receive the SQL query from the client.
 - Parse the SQL query check for syntax errors.
 - Check if the user has required privilege against the objects used in the SQL query.
 - Check if the objects used in the SQL actually exists.
 - Prepare the execution plan to execute the SQL query and pass it to BYNET.
 - Receives the results from the AMPs and send to the client.

Components of Teradata (cont.)

- **Message Passing Layer:** Message Passing Layer called **BYNET**, is the networking layer in Teradata system. It allows the communication between PE and AMP and also between the nodes.
- It receives the execution plan from Parsing Engine and sends to AMP.
- Similarly, it receives the results from the AMPs and sends to Parsing Engine.

Components of Teradata (cont.)

- **Access Module Processor (AMP):** AMPs, called as Virtual Processors (vprocs) are the one that actually stores and retrieves the data.
- AMPs receive the data and execution plan from Parsing Engine, performs any data type conversion, aggregation, filter, sorting and stores the data in the disks associated with them.
- Records from the tables are evenly distributed among the AMPs in the system.
- Each AMP is associated with a set of disks on which data is stored. Only that AMP can read/write data from the disks.

Storage Architecture


- When the client runs queries to insert records, Parsing Engine sends the records to BYNET.
- BYNET retrieves the records and sends the row to the target AMP.
- The AMP stores these records on its disks

Retrieval Architecture

- When the client runs queries to retrieve records, Parsing Engine sends a request to BYNET.
- BYNET sends the retrieval request to appropriate AMPs.
- Then AMPs search their disks in parallel and identify the required records and sends to BYNET.
- BYNET then sends the records to Parsing Engine which in turn will send to the client.

How does Teradata store rows?

- Random, automatic data distribution & placement
- Real-time, automatic data reorganization



*With Teradata there is no sense of ORDER,
therefore there is no sense of DISORDER,
eliminating the need to REORDER!*

Quick reminder

- **Primary Key:** Primary key is used to uniquely identify a row in a table.
 - No duplicate values are allowed in a primary key column and they cannot accept NULL values.
 - It is a mandatory field in a table.
- **Foreign Key:** Foreign keys are used to build a relationship between the tables.
 - A foreign key in a child table is defined as the primary key in the parent table.
 - A table can have more than one foreign key.
 - It can accept duplicate values and also null values.
 - Foreign keys are optional in a table.

Primary Indexes

- The mechanism used to assign a row to an AMP.
- A table ~~must~~ have a **Primary Index** that cannot be changed.
 - From Teradata 13.00 tables may *not* have a primary index. Rows are randomly distributed to AMPs.
 - No PI tables are typically used as staging tables for initial load by FastLoad or TPump Array Inserts (because load is faster).
- Primary Index can be unique (**UPI**) or non-unique (**NUPI**)

Quick quiz:


- Consider a transaction table with columns `OrderNumber` , `CustomerId` , `OrderDate` and `Total` .
- How would you assign primary keys and primary indices to make your life happier?
- **Happiness** is defined as fast running queries.

Solution

- It depends.
- Rows can be distributed using a `UPI` (in this case, `OrderNumber` , which is also the PK) or a `NUPI` (in this case, `CustomerId` or `OrderDate`).
- In the first case, the distribution of the rows is non-skewed across AMPs, while in the second case we will have a less even row distribution.
- But the catch is on **what level of granularity** do you need for analysis.


Why this matters?: JOIN

- In order to process records in a JOIN, they have to be on the same processing unit
- Most RDBMS have to send the records to a single processing unit, to perform the JOIN



Avoiding bottlenecks

- The `JOIN` is performed balanced on all nodes if the primary index is chosen properly.
- Each node performs a smaller part of the `JOIN`.
- Similar bottlenecks (and solutions) apply for `GROUP BY` and `ORDER` queries.



Hashing algorithm

- A row is assigned to a particular AMP based on the primary index value. Teradata uses a hashing algorithm to determine which AMP gets the row.
- How this looks like:
 - The client submits a query.
 - The parser receives the query and passes the PI value of the record to the hashing algorithm.
 - The hashing algorithm hashes the primary index value and returns a 32 bit number, called **Row Hash**.

Hashing algorithm (cont.)

- The higher order bits of the row hash (first 16 bits) are used to identify the hash map entry.
- The hash map is an array of buckets which contains specific AMP number.
- BYNET sends the data to the identified AMP.
- AMP uses the 32 bit Row hash to locate the row within its disk.
- If there is any record with same row hash, then it increments the **uniqueness ID**.
- The pair of Row hash and uniqueness ID is called **Row ID**.
- Tables are sorted by their Row ID.

Setting up the environment

Pre-requisites

To get started, you need to download:


- Teradata Express 16.20 VM Image.
- VMWare Workstation 15 Player.
- Teradata tools and utilities.
- JDK 8 and other dependencies (see the download page for TD Express).
- You can download these from: [Teradata Downloads page](#).
- An account is required, you can create one for free.

Launch your VM

- In VMWare Player's main window, right click on the VM name and select "Virtual Machine Settings".
- In the settings window, click on "Network Adapter" and among the options on the right, select "Host-only: a private network shared with the host".
- Click ok and boot your virtual machine.
- Login credentials are root/root

Connect to your DB

- Open a terminal and type `ifconfig`. You should see an `inet addr` value as below.



The screenshot shows a VMware Workstation window titled "TDEExpress16.20_Sles11 - VMware Workstation ...". Inside the window is a terminal window titled "Terminal". The terminal displays the output of the `ifconfig` command. The output shows two network interfaces: `eth0` and `lo`. The `eth0` interface has an `inet addr` of `192.168.47.129`. The `lo` interface has an `inet addr` of `127.0.0.1`.


```
screen ("TERADATA EXPRESS") EVALUATION AND DEVELOPMENT
LICENSE AGREEMENT), including the restriction that this
evaluation copy is not for production use.
TDEExpress1620_Sles11:~ # ifconfig
eth0      Link encap:Ethernet  HWaddr 00:0C:29:60:94:29
          inet addr:192.168.47.129  Bcast:192.168.47.255  Mask:255.255.255.0
          inet6 addr: fe80::20c:29ff:fe60:9429/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:23 errors:0 dropped:0 overruns:0 frame:0
          TX packets:15 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:3424 (3.3 Kb)  TX bytes:1572 (1.5 Kb)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          inet6 addr: ::1/128 Scope:Host
          UP LOOPBACK RUNNING  MTU:16436  Metric:1
          RX packets:839 errors:0 dropped:0 overruns:0 frame:0
          TX packets:839 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:86289 (84.2 Kb)  TX bytes:86289 (84.2 Kb)

TDEExpress1620_Sles11:~ #
```

Test connection

- On your PowerShell or UNIX console, ping the IP address of your DB with the command `ping 192.168.47.129`

A screenshot of a Windows PowerShell terminal window. The window title is "Windows PowerShell" with standard minimize, maximize, and close buttons. The terminal text shows the command "ping 192.168.47.129" being executed. The output displays four successful replies from the target IP address, each with 32 bytes of data, a time of 1ms or less, and a TTL of 64. Below the replies, the ping statistics are shown, indicating 4 packets sent, 4 received, and 0% loss, with round trip times of 0ms minimum, 1ms maximum, and 0ms average. The prompt "PS C:\Users\HP>" is visible at the bottom.

```
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

PS C:\Users\HP> ping 192.168.47.129

Pinging 192.168.47.129 with 32 bytes of data:
Reply from 192.168.47.129: bytes=32 time=1ms TTL=64
Reply from 192.168.47.129: bytes=32 time<1ms TTL=64
Reply from 192.168.47.129: bytes=32 time<1ms TTL=64
Reply from 192.168.47.129: bytes=32 time<1ms TTL=64

Ping statistics for 192.168.47.129:
    Packets: Sent = 4, Received = 4, Lost = 0 (0% loss),
    Approximate round trip times in milli-seconds:
        Minimum = 0ms, Maximum = 1ms, Average = 0ms
PS C:\Users\HP>
```

Connect with SQLA

- Now open SQL Assistant, and create a new connection using the IP address of your DB.
- Default credentials are dbc/dbc.
- Run a simple query.

The screenshot displays the Teradata SQL Assistant (teradata-poland) application. The interface includes a menu bar (File, Edit, View, Tools, Window, Help), a toolbar, and a Database Explorer on the left showing the 'teradata-poland (Teradata)' connection with a 'DBC' folder expanded. The main query editor contains the SQL statement: `select * from dbc.dbcinfo`. Below the editor, the 'Answerset 1' pane shows the results of the query in a table format. The 'History' pane at the bottom lists the executed queries, including the current one.

InfoKey	InfoData
1 LANGUAGE SUPPORT MODE	Standard
2 RELEASE	16.20.23.01
3 VERSION	16.20.23.01

Date / Time	Source	Elapsed	Rows	Result	Notes	SQL Statement	Length	Stmts
5/8/2019 10:10:40	teradata-poland	00:00:00	3			/**/	33	1
5/6/2019 11:22:31	tutorial wroclaw	00:00:00	1			select sum(maxperm), sum(currentperm) from dbc.diskspace	56	1
5/6/2019 11:22:00	tutorial wroclaw	00:00:01	60			select * from dbc.diskspace	27	1
5/6/2019 11:19:59	tutorial wroclaw	00:00:00	1			SELECT Hashamp()	16	1
5/6/2019 11:18:42	tutorial wroclaw	00:00:00	1			SELECT CURRENT_DATE	19	1
5/6/2019 11:18:39	tutorial wroclaw	00:00:00	1			SELECT CURRENT_TIME	19	1
5/6/2019 11:18:08	tutorial wroclaw	00:00:00	1			SELECT CURRENT_DATE	19	1

teradata-poland Line 3 100% 10:52

SQLA Shortcuts

- **F2**: It will open query builder, with syntax for all SQL queries.
- **F5**: Execute SQL query.
- **F6**: Explain plan for SQL query.
- **F9**: Execute SQL queries in parallel.
- **F10**: Abort SQL query.
- **F11**: Display last error encountered.
- **Ctrl + N**: New SQL query window.
- **Ctrl + Q**: Format SQL query.
- **Ctrl + U**: Convert to UPPERCASE.
- **Ctrl + H**: Find and replace.

More inside SQLA

- Go to **Tools | Options | Query** to change the behavior if needed.
- **Query Builder.**
- **Beyond SQLA:** Teradata Studio / Teradata Studio Express.