

Performance Monitoring

Top 5 reasons for slow queries

1. Bad Query

- This is the most common reason for long running query.
- Your query is not using any kind of index: Primary Index, Secondary Index or Partition Primary Index.
- The tables are getting redistributed and those big tables taking long time in it.

2. Delay Time

- Check for Delay time in DBQL (`dbc.dbqlogtbl`) for your query.
- It may happen that query execution time is in seconds however it was in Delay queue for hours.

3. Blocking

- It may happen that your query is blocked by some other query.
- Also watch out for TDM jobs. Such jobs generally apply exclusive locks on the table, thereby blocking all other queries requesting the same table.
- You can check this Teradata Viewpoint Monitor Portlet.

4. Server State

- Generally Teradata Server status is classified as Healthy, Degraded, Critical and Down.
- As the number of concurrent users increase the load on server also increases.
- With more load on server and AMP doing more work, your query may take more time than usual time.

5. Skewness

- Check skewness factor and change the PI if needed.

Improving your query

DISTINCT

- `DISTINCT` is better for columns with a low number of rows per value:
Number of rows < Number of AMPs
- `GROUP BY` is better for columns with a large number of rows per value:
Number of rows > Number of AMPs

Date Comparison

- When comparing values of date in a particular range, the query may result in product join.
- This can be avoided with the usage of `SYS_CALENDAR.CALENDAR`, which is Teradata's in-built database.

Example

```
select
t2.a1,t2.a2,t2.a3,t2.a4
from
table2 t2
join table3 t3
on t2.a1=t3.a1
and t2.a5_dt>=t3.a4_dt
and t2.a5_dt<=t3.a5_dt;
```

```
select
t2.a1,t2.a2,t2.a3,t2.a4
from table2 t2
join SYS_CALENDAR.CALENDAR sys_cal
on sys_cal.calendar_date = t2.a5_dt
join table3 t3
on t2.a1=t3.a1
and sys_cal.calendar_date >=t3.a4_dt
and sys_cal.calendar_date <=t3.a5_dt;
```

Identify suspect queries

Beyond EXPLAIN

- **Product Join Indicator:** the ratio of CPU Seconds to IO for a query. $(\text{AMPCPUTime} * 1000) / \text{TotalIOCount}$
 - ≥ 3 : the query should be reviewed.
 - ≥ 6 : potentially an unnecessary product join.

Beyond EXPLAIN (cont.)

- **Unnecessary IO Indicator:** is the ratio of IO to CPU Seconds,
$$\text{TotalIOCount} / (\text{AMPCPUTime} * 1000)$$
 - Number of rows read / number of rows processed.
 - ≥ 3 : the query should be reviewed to eliminate full-table scans and possibly redistribution steps.
 - Good indicator

Beyond EXPLAIN (cont.)

- Product Join Indicator, $(\text{AMPCPUTime} * 1000) / \text{TotalIOCount}$
- PJI is the measure of how CPU intensive your query is.
 - If $\text{PJI} \geq 3$ you should review the query, $\text{PJI} \geq 6$ it might be a product join.
 - This value is high during a product join, but not only.

Beyond `EXPLAIN` (cont.)

- Both metrics are available in Viewpoint's Query Monitor portlet, and every individual query has these values displayed when you click on the session ID.
- You can also track them using `dbqlogtbl` & `dbqsqltbl` tables in `dbc` .

Set up dbq1

- You can check if this is activated:

```
SELECT * FROM dbc.dbq1rulesv;
```

- If not activated, it can be set up by the DBA in BTEQ .

```
begin query logging with objects,  
sql, usecount, utilityinfo LIMIT SQLTEXT=0 on all;
```

```
begin query logging with objects,  
sql limit threshold = 5 elapsedsec  
and sqltext=0 on VIEWPOINT;
```

Monitoring

```
SET QUERYBAND = 'Version=1' FOR SESSION;

SELECT
AMPCPUTIME,
(FIRSTRESPTIME-STARTTIME DAY(2) TO SECOND(6)) RUNTIME,
SPOOLUSAGE/1024**3 AS SPOOL_IN_GB,
AMPCPUTIME*1000/TOTALIOCOUNT AS PJI,
TOTALIOCOUNT/(AMPCPUTIME*1000) AS UIOI
FROM DBC.DBQLOGTBL
WHERE QUERYBAND = 'Version=1';
```

Monitoring (cont.)

The query will return:

- Total CPU Usage
- Spool Space needed
- PJI/LHR (ratio between CPU and IO usage)
- CPU Skew
- Skew Impact on the CPU
- **Goal:** cut total CPU usage, consumed spool space and skew.

Prevention

Good Schema Design

- In a perfect world, we don't need to optimize our queries because the schema is just great.
- Data is **normalized**: tables are organized to reduce disk space *and* redundancy.
- We want to avoid situations like having `EmployeeSalary` in a `Customers` table. Or, even worse, a single, huge table with everything.
- More rows, less columns.

Normalization

- Different rules for normalization.
- Depending how many of those are followed, we say that the database is in First Normal Form, Second Normal Form (5 and counting).
- Third Normal Form (3NF) is the standard recommended by Teradata, and usually the last necessary.

HOW STANDARDS PROLIFERATE:
(SEE: A/C CHARGERS, CHARACTER ENCODINGS, INSTANT MESSAGING, ETC.)



First Normal Form

- Eliminate columns with repeated information in individual tables.
- Create a separate table for each set of related data.
- Identify each set of related data with a primary key.

Second Normal Form

- Create separate tables for sets of values that apply to multiple records.
- Relate these tables with a foreign key.
- **Example:** `CustomerAddress` . You need this information in your `Orders` , `Invoices` and `Customers` table. Instead of duplicating it, it should be only on `Customers` table.

Third Normal Form

- Eliminate fields that do not depend on the key.
- In the previous example, that could mean storing `CustomerAddress` in its own `Addresses` table.
- Theoretical advantages, but maybe not practical: it would mean having a lot of small tables in some cases, which can degrade performance.

Normalizing an Example Table

These steps demonstrate the process of normalizing a fictitious student table.

Unnormalized table:

StudentId	Advisor Adv-Room	Class1	Class2	Class3
1022	Jones	412	101-07	143-01
4123	Smith	216	201-01	11-02

First Normal Form: No duplicate columns

StudentId	Advisor	Adv-Room	ClassId
1022	Jones	412	101-07
1022	Jones	412	143-01
1022	Jones	412	159-02
4123	Smith	216	201-01
4123	Smith	216	211-02
4123	Smith	216	214-01

Second Normal Form: Separate tables

StudentId	Advisor	Adv-Room
1022	Jones	412
4123	Smith	216

StudentId	ClassId
1022	101-07
1022	143-01
1022	159-02
4123	201-01
4123	211-02
4123	214-01

Third Normal Form

StudentId	Advisor
1022	Jones
4123	Smith

Name	Room	Dept
Jones	412	42
Smith	216	42

Why not to normalize?

- Joins are expensive.
- Prototyping should be *quick and dirty*.
- If using NoSQL databases.

Other tips

- Primary key *need not be* a business attribute (e.g. ID Number).
These may change!
- Multi-column Primary Key = Bad idea in general.

Other tips (cont.)

- Add constraints on the database, specially unique (when relevant).
- **Avoid bad names,**
 - `id` in `Customers` and `product_id` in `Products` .
 - Non-ASCII characters, whitespace, reserved words.
 - Mixing plural and singular.
 - Or mixing underscore and CamelCase
- Hard Delete vs Soft Delete: Flag as 1/0 for active/inactive instead of deleting.