

Performance monitoring

Top 5 reasons for slow queries

1. Bad Query

- This is the most common reason for long running query. Your query is not using any kind of index: Primary Index, Secondary Index or Partition Primary Index. The tables are getting redistributed and those big tables taking long time in it.

2. Delay Time

- Check for Delay time in DBQL for your query. It may happen that query execution time is in seconds however it was in Delay queue for hours.

3. Blocking

- It may happen that your query is blocked by some other query.
- Also watch out for TDM jobs. Such jobs generally apply Exclusive locks on the table thereby blocking all other queries requesting the same table.
- You can check this Teradata Viewpoint Monitor Portlet.

4. Server State

- Generally Teradata Server status is classified as Healthy, Degraded, Critical and Down.
- As the number of concurrent users increase the load on server also increases.
- With more load on server and AMP doing more work, your query may take more time than usual time.

5. Skewness

- Check skewness factor and change the PI if needed.

Improving your query

DISTINCT

- `DISTINCT` is better for columns with a low number of rows per value:

Number of rows < Number of AMPs

- `GROUP BY` is better for columns with a large number of rows per value:

Number of rows > Number of AMPs

IN

- No theoretical limit to the number of items contained within an `IN` clause.
- **But:** performance degradations if the list starts to exceed a few hundred items.
- Optimizer tries to build a spool file when a large `IN` list is encountered.
- When the number of values exceeds the acceptable limit, the hard-coded values can be inserted into a volatile table and the `IN` clause can be made to an join.

Date Comparison

- When comparing values of date in a particular range, the query may result in product join. This can be avoided with the usage of `SYS_CALENDAR.CALENDAR` , which is Teradata's in-built database.

Example

```
select
t2.a1, t2.a2, t2.a3, t2.a4
from
table2 t2
join table3 t3
on t2.a1=t3.a1
and t2.a5_dt>=t3.a4_dt
and t2.a5_dt<=t3.a5_dt;
```

```
select
t2.a1, t2.a2, t2.a3, t2.a4
from table2 t2
join SYS_CALENDAR.CALENDAR sys_cal
on sys_cal.calendar_date = t2.a5_dt
join table3 t3
on t2.a1=t3.a1
and sys_cal.calendar_date >=t3.a4_dt
and sys_cal.calendar_date <=t3.a5_dt;
```

Identify suspect queries

Beyond **EXPLAIN**

- **Product Join Indicator:** the ratio of CPU Seconds to IO for a query. $(\text{AMPCPUTime} * 1000) / \text{TotalIOCount}$
 - ≥ 3 : the query should be reviewed.
 - ≥ 6 : potentially an unnecessary product join.

Beyond `EXPLAIN` (cont.)

- **Unnecessary IO Indicator:** is the ratio of IO to CPU Seconds,
$$\text{TotalIOCount} / (\text{AMPCPUTime} * 1000)$$
 - ≥ 3 : the query should be reviewed to eliminate full-table scans and possibly redistribution steps.
 - UIOI is a reasonable indicator to identify queries that may benefit from additional statistics to indexing improvements.
- Requires `dbqlogtbl` & `dbqlsqltbl` tables in `dbc`.

Set up dbq1

- You can check if this is activated:

```
SELECT * FROM dbc.dbq1rulesV;
```

- If not activated, it can be set up by the DBA in BTEQ .

```
begin query logging with objects,  
sql, usecount, utilityinfo LIMIT SQLTEXT=0 on all;
```

```
begin query logging with objects,  
sql limit threshold = 5 elapsedsec  
and sqltext=0 on VIEWPOINT;
```


Beyond **EXPLAIN** (cont.)

- PJI is the measure of how CPU intensive your query is.
 - If PJI is relatively high for a query, then the query takes many CPU cycles for the given number of I/Os.
 - This value is high during a product join, but not only.

Beyond **EXPLAIN** (cont.)

- If UII is relatively high for a query, then it could mean that many I/O blocks are read, meaning how IO intensive your query is, but a relatively low number of rows is actually processed.
- If it is a full table scan with only a few rows qualifying, then an index could reduce the I/O consumption in this case.
- Both metrics are available in Viewpoint's Query Monitor portlet, and every individual query has these values displayed when you click on the session ID.

Using QUERYBAND

```
SET QUERYBAND = 'Version=1;' FOR SESSION;  
SELECT  
  AMPCPUTIME,  
  SPOOLUSAGE/1024**3 AS SPOOL_IN_GB,  
  100-100*((AMPCPUTIME/(HASHAMP()+1))  
    /NULLIFZERO(AMPCPUTIME)) AS CPU_SKEW,  
  MAXAMPCPUTIME*(HASHAMP()+1) AS CPU_IMPACT,  
  AMPCPUTIME*1000/(TOTALIOCOUNT) AS LHR  
FROM  
  DBC.DBQLOGTBL  
WHERE QUERYBAND = 'Version=1;'
```

- <https://www.dwhpro.com/teradata-sql-tuning-top-10/>

Using QUERYBAND (cont.)

The query will return:

- Total CPU Usage
- Spool Space needed
- LHR (ratio between CPU and IO usage)
- CPU Skew
- Skew Impact on the CPU
- **Goal:** cut total CPU usage, consumed spool space and skew.