Indices (or Indexes)

In this lecture

- Types of indexes in Teradata.

- Primary indexes.

- Partitioned primary indexes.

- Join indexes.

Motivation: Why do we care?

- Indexes are the **most important part** of designing the database structure.

- Indexes not only provide an effective way to store data, but also help in determining effective access paths to data.

Indexes in Teradata

- Unique/non-unique primary index (**UPI/NUPI**)

- Partitioned primary index (**PPI**)

- Unique/non-unique secondary index (**USI/NUSI**)

- Partitioned primary index (**PPI**)

- Join index (**JI**)

Index families

- **Primary indexes**: To distribute and retrieve data rows in a table. Storage and maintenance are free.

- **Partitioned primary index**: A table organization to optimize the physical database design for range constraint queries. Storage is 2 bytes per row.

- **Raw data extensions**: Any structure that duplicates or points to primary data for purposes of better performance. These are like secondary indexes (SI) or join indexes (JI). Storage and maintenance are **not free**.

Primary Indexes

- Cannot be modified once the table is non-empty...

- but you can either:
  - create a new table with the index structure you need.
  - copy the data out to another table (volatile or permanent), modify the index and move the data back.

Primary Indexes (cont.)

- You can modify the index of an empty table:

```
ALTER TABLE <table> MODIFY
PRIMARY INDEX Index_Name(col1, col2, ...)
```

- Should you modify primary indexes? That depends.
- Usually other workarounds are possible, since PI determines the location of the data, so re-indexing involves moving data around (often unwanted).
- Workarounds: partitioned indexes or secondary indexes.

Choosing a good primary index

- A good primary index should satisfy the following three properties:
  - Access
  - Distribution
  - Volatility

Choosing a good primary index (cont.)

- **Access**: Choose the column that provides the best access path to the data.
  - Are the individual records commonly selected? Then use the PK.
  - Do you typically join this table? Then consider rather the join columns.

Choosing a good primary index (cont.)

- **Distribution**: Help Teradata distribute data evenly choosing columns that would have a regular distribution. Avoid implementing numeric values as `VARCHAR/CHAR` to avoid hash collisions.

Choosing a good primary index (cont.)

- **Volatility**: Choose a column with stable data values.
- This helps to reduce overhead of data maintenance (moving data around).

Example: Choosing a good index

- Suppose you have the following tables:

- `Order(PRIMARY KEY(OrderNumber))`.

- `LineItem(PRIMARY KEY(OrderNumber, ItemNumber))`.

Example: Choosing a good index (cont.)

- Orders are commonly looked up by `OrderNumber` in the `Order` table.
- Line items are typically accessed by joining the `Order` table to the `LineItem` table on `OrderNumber`.

Example: Choosing a good index (cont.)

- `OrderNumber` s are unique in the `Order` table, and non-unique in the `LineItem` table.
- Since `OrderNumber` is the key to the `Order` table `OrderNumber` values do not typically change.
- **Question:** What should be the primary indices on each table, and why?
- **Hint:** think in terms of access, distribution and volatility.

Example (cont.)

- `OrderNumber` is the only choice for PI in `Order` , and it satisfies access, distribution, volatility.

- The PK of `LineItem` satisfies distribution and volatility, but not access: the table has to be re-distributed every time it is joined with `Order` !

- `OrderNumber` satisfies the three conditions in `LineItem` , hence this should be selected as PI of both tables.

Partitioned primary index

Partitioned primary index

- With a PPI, rows are sent to different AMPs, but also *local partitions* within each AMP are created.

- Normal PI access remains unchanged, but in the case of a range query, each AMP is able to localize the search on specific partitions within its workspace.

- This means that Teradata Optimizer knows the portions of a range of values stored and scan only those parts in the table.

Example

- `03 Indexes.sql`
- Create `salary_non_ppi` and run `EXPLAIN` on the query:
  - `SELECT * from salary_non_ppi where dob <= '2017-09-01';`
- **Red flags:**
  - `all-rows scan` : polite synonym for **really slow**.
  - `no confidence` : Optimizer has not collected statistics on this table. Usually a bad sign, this is the least of our problems right now.

Example (cont.)

- So how do we get rid of that?

- Partition by range! (DDL for `salary_ppi`).

- Four different types of PPI.

Case partitioning

```
/*CASE partition*/
CREATE TABLE SALES_CASEPPI
(
  ORDER_ID INTEGER,
  CUST_ID INTERGER,
  ORDER_DT DATE,
)
PRIMARY INDEX(ORDER_ID)
PARTITION BY CASE_N(ORDER_ID < 101,
ORDER_ ID < 201,
ORDER_ID < 501,
NO CASE,UNKNOWN);
```

Range-based partitioning

```
/*Range Partition table*/
CREATE volatile TABLE EMP_SAL_PPI
(
id INT,
Sal int,
dob date,
bonus int
) primary index( id)
PARTITION BY RANGE_N (dob BETWEEN DATE '2017-01-01'
AND DATE '2017-12-01' EACH INTERVAL '1' DAY)
on commit preserve rows;
```

Multi-level partitioning

```sql
CREATE TABLE SALES_MLPPI_TABLE
(
ORDER_ID INTEGER NOT NULL,
CUST_ID INTERGER,
ORDER_DT DATE,
)
PRIMARY INDEX(ORDER_ID)
PARTITION BY (
        RANGE_N(
                ORDER_DT BETWEEN DATE '2017-08-01'
                AND DATE '2017-12-31'
            EACH INTERVAL '1' DAY)

    CASE_N (ORDER_ID < 1001,
        ORDER_ID < 2001,
        ORDER_ID < 3001,
        NO CASE, UNKNOWN));
```

Character-based partitioning:

```sql
/*CHAR Partition*/
CREATE TABLE SALES_CHAR_PPI (
ORDR_ID INTEGER,
EMP_NAME VARCHAR (30) CHARACTER,
PRIMARY INDEX (ORDR_ID)
PARTITION BY CASE_N (
EMP_NAME LIKE 'A%', EMP_NAME LIKE 'B%',
EMP_NAME LIKE 'C%', EMP_NAME LIKE 'D%',
EMP_NAME LIKE 'E%', EMP_NAME LIKE 'F%',
NO CASE, UNKNOWN);
```

Secondary Index

Secondary Index

- A table can contain **only one** primary index.

- More often, you will come across scenarios where the table contains other columns, using which the data is frequently accessed.

- Secondary indexes are used to avoid full table scan in those cases.
  - Optional and not involved in data distribution.
  - Stored in sub tables. These tables are built in all AMPs.
  - They can be created during table creation or after a table is created.
  - They also require maintenance since the **sub-tables need to be updated for each new row.**

The truth...

- Teradata runs extremely well **without** secondary indexes.

- Only recommended when queries that are run over and over.

- If the tables are modified, need to be recreated! Maintenance overhead.

```
CREATE UNIQUE INDEX (Column/Columns) ON <tablename >;

CREATE INDEX (Column/Columns) ON <tablename >;
```

Join Indexes

Join Index

- A `JOIN INDEX` is a materialized view. Its definition is permanently stored and the data is updated whenever the base tables referred in the join index is updated.

- JOIN INDEX may contain one or more tables and also contain pre-aggregated data. Join indexes are mainly used for improving the performance.
    - Different types of join indexes available.
    - Single Table Join Index (STJI)
    - Multi Table Join Index (MTJI)
    - Aggregated Join Index (AJI)

Example

Suppose we have the following tables:

```
CREATE SET TABLE EMPLOYEE,FALLBACK
(
EmployeeNo INTEGER,
FirstName VARCHAR(30) ,
LastName VARCHAR(30) ,
DOB DATE FORMAT 'YYYY-MM-DD',
JoinedDate DATE FORMAT 'YYYY-MM-DD',
DepartmentNo BYTEINT
)
UNIQUE PRIMARY INDEX ( EmployeeNo );
```

Example (cont.)

```
CREATE SET TABLE SALARY,FALLBACK
(
EmployeeNo INTEGER,
Gross INTEGER,
Deduction INTEGER,
NetPay INTEGER
)
PRIMARY INDEX ( EmployeeNo )
UNIQUE INDEX (EmployeeNo);
```

STJI: Example

Let's create a JOIN index on the `Employee` table.

```
CREATE JOIN INDEX Employee_JI
AS
SELECT EmployeeNo,FirstName,LastName,
BirthDate,JoinedDate,DepartmentNo
FROM Employee
PRIMARY INDEX(FirstName);
```

- When the user submits a query with a `WHERE` clause on `EmployeeNo` then the UPI is used.

- If the query is on `FirstName`, then the system may access it on `Employee_JI`.

- On other columns (e.g. `LastName`, then full table scan is necessary.

MTJI: Example

```
CREATE JOIN INDEX Employee_Salary_JI
AS
SELECT a.EmployeeNo,a.FirstName,a.LastName,
a.BirthDate,a.JoinedDate, a.DepartmentNo
,b.Gross,b.Deduction,b.NetPay
FROM Employee a
INNER JOIN Salary b
ON(a.EmployeeNo=b.EmployeeNo)
PRIMARY INDEX(FirstName);
```

- If you run a query joining these tables, then Optimizer may choose to access the data from the join index directly.

- You can verify what will happen with `EXPLAIN`.

AJI: Example

```
CREATE JOIN INDEX Employee_Salary_JI
AS
SELECT a.DepartmentNo,SUM(b.NetPay) AS TotalPay
FROM Employee a
INNER JOIN Salary b
ON(a.EmployeeNo=b.EmployeeNo)
GROUP BY a.DepartmentNo
Primary Index(DepartmentNo);
```

- If a table is consistently aggregated on certain columns, you can create an aggregate join index.

- It supports only `SUM` and `COUNT`.